




(a) Decision profiles




(b) Input length distribution




(e) Evaluation taxonomy



(c) Query-text deduplication retention



(d) Reference-label balance



(a) Shared 0-100 accuracy scale. (b) Reference token counts, not native cost. (c) 128 query-only samples per task; threshold sensitivity, not CI. (d) Full frozen reference labels. (e) Angles count taxonomy leaves, not sample shares. Full intent tests; historical tasks retain frozen subsets.