




(a) Decision profiles




(b) Input length distribution




(e) Evaluation taxonomy



(c) Query-text deduplication retention



(d) Reference-label balance




(a) Shared 0-100 accuracy scale. (b) Reference token counts, not native cost. (c) 128 query-only samples per task; threshold sensitivity, not CI. (d) Full frozen reference labels. (e) Angles count taxonomy leaves, not sample shares. Full intent tests; historical tasks retain frozen subsets.

(a) Decision profiles



Six representative models across the same eight domains.




Radial length = domain accuracy, 0-100; every sector has the same scale.

Point estimates only; no significance claim. Full 17-model values are supplied separately.

(b) Input length distribution

Common tokenizer; every candidate is retained.




Log-scaled x axis; densities normalized within each group. Gaussian bandwidth = 0.12 log10 units.

Compact gold-free state + target question JSON, Qwen3.8 tokenizer; not native/server token cost.

(c) Query-text deduplication retention

Question text only; not whole-dataset diversity.



128 hash-selected queries per task. MiniLM cosine threshold = 0.80; deterministic greedy retention.


Lines show threshold sensitivity (0.75-0.85), NOT confidence intervals. Query-only proxy, not difficulty.

ContractNLI reuses hypothesis sentences; evidence is excluded, so this is NOT whole-corpus diversity.

(d) Reference-label balance



Declared label space, including labels with zero support.




$100 \times H(\text{gold labels}) / \log(K)$, using every original scored case and the complete K-choice ontology.

Balance is not accuracy. When2Call has 3/4 observed labels; CLINC includes 1,000 OOS examples.

(e) Evaluation taxonomy



Eleven classification tasks, with separate diagnostic tracks.



Angular width counts terminal taxonomy leaves, NOT examples, score weights, or model coverage.
Only classification enters the eight-domain index. Diagnostics stay separate; coverage differs by track.