


Equal scores can conceal different mistakes

Kev-27B v2 and Qwen3.8-27B, paired on the same original questions



Numbers inside segments are case counts. Scientific base accuracy is 285/339 for each model, yet each uniquely solves 21 cases.

The right column is a retrospective best-of-two upper bound, not a tested selector or deployable routing result.