


Same-size contrasts are not the same experiment

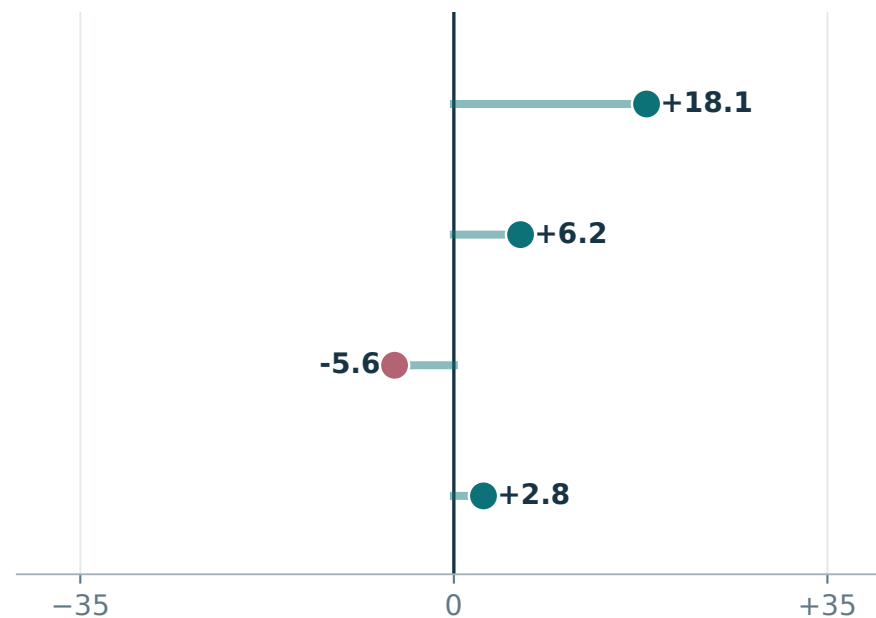
Each dot is a paired-sample difference in the original-condition score; positive values favor Kev.

Kev-0.8/4/9B vs Qwen3.5; Kev-27B vs Qwen3.8. All pairs share identical requests, but not a common training history.


EXECUTABLE POLICY



CONTRACT LABELS



SCIENTIFIC EVIDENCE



* Kev-9B is the earlier pinned checkpoint. † Kev-27B v2 shares an already post-trained Qwen3.8-27B base.

Descriptive differences only: no causal claim about scale, pretraining, or decision-model architecture.