


Stable total accuracy hides a 60-case answer migration

Qwen3 8B · 180 cited claim–abstract pairs · 60 references per class



Overall correct: 104 → 101 / 180 (-1.7 pp; paired 95% CI [-7.8, +4.4])

60 / 180 reversed predictions changed, all from base NOINFO; exact repeat changed 0 / 180.

NOINFO: explicitly cited paper without annotated abstract evidence (130 candidates, 60 selected).
The LLM adapter also changes answer-code mapping on reversal; this is interface sensitivity, not pure semantics.