



Default distributions versus exact event probabilities.

96 exact-probability cases, 48 paired settings. Not single-outcome accuracy, not correctness-confidence calibration; no fitting on test probabilities.



Lower is better · $\text{sum}(\text{predicted} - \text{exact})^2$
Reference probabilities are checked as exact fractions. Error bars: 95% paired-setting bootstrap intervals. Invalid outputs are reported with valid denominators.

This small diagnostic comes from the Intern repository; it is not an independent broad probability benchmark or evidence of general decision competence.